

# Report Of NLP Semantic Matching Project Report

## *Intelligent Equipment Retrieval System Using Hybrid Semantic & Lexical Search*

**Team Members:** Sarra Dhouaifi & Tharaa Oueslati

**Course:** DevOps

**Date:** January 2026

**Project Repository:**

<https://github.com/dhouaifisarra/NLP-Semantic-Matching-Project>

---

## 1. Project Overview

This project focuses on building an **Equipment Semantic Matching System** capable of matching free-text user queries to the most relevant equipment items from a catalog. The system combines **semantic search, lexical similarity, and a hybrid scoring strategy** to achieve robust and accurate matching results.

The solution is delivered as:

- A **FastAPI-based REST API** for programmatic access and evaluation
  - A **Web-based User Interface** for manual testing and feedback collection
  - A **Dockerized application** to ensure reproducibility and ease of deployment
- 

## 2. Problem Statement

In many industrial and technical domains, users search for equipment using **natural language queries** that may not exactly match catalog item names. Traditional keyword-based search systems often fail to capture semantic similarity, leading to poor retrieval performance.

The objective of this project is to design and implement a system that:

- Understands the **semantic meaning** of user queries
- Retrieves the most relevant equipment even when wording differs

- Provides measurable evaluation metrics
  - Allows both automated and manual evaluation
- 

### 3. Dataset Description

#### 3.1 Synthetic Data Generation

Due to the absence of publicly available labeled datasets for equipment-query matching, this project relies on **synthetic data generated using an AI-based text generation approach**.

The synthetic dataset was designed to:

- Simulate realistic industrial equipment names
- Generate diverse user search queries
- Avoid the use of real, sensitive, or proprietary data

Using synthetic data ensures **reproducibility, ethical compliance, and full control** over the dataset structure.

#### 3.2 Data Files

- `equipment_catalog.csv`: Raw synthetic equipment catalog
  - `equipment_clean.csv`: Preprocessed and cleaned equipment data
  - `test_queries.csv`: Synthetic test queries for evaluation
  - `faiss_index.idx`: Semantic embeddings index
  - `faiss_index.idx_mapping.pkl`: ID → equipment name
- 

### 4. Data Preprocessing

Before indexing and retrieval, the data undergoes several preprocessing steps:

- Text normalization (lowercasing)
- Removal of special characters and accents
- Deduplication of equipment entries
- Removal of punctuation and extra whitespace
- Construction of a clean text field used for embeddings

This preprocessing ensures consistency between catalog entries and user queries.

---

## 5. Methodology

The core goal of this project is to retrieve the most relevant equipment items from a catalog given a query. We combine **semantic search**, **lexical search**, and a **hybrid approach**, with **score calibration** for confidence estimation.

Below is a detailed explanation of each step and algorithm.

### 5.1 Semantic Search

#### Purpose:

Semantic search captures the meaning of a query rather than just exact word matches. This helps retrieve relevant items even when the user uses synonyms or paraphrases of equipment names.

#### Algorithm & Steps:

##### 1. Embedding Generation

- We use the pre-trained sentence-transformer/all-MiniLM-L6-v2 model from HuggingFace.
  - This model transforms text (equipment names and queries) into dense vector embeddings in a high-dimensional space.
  - Each equipment name is embedded into a vector of size 384 (the default dimension of MiniLM-L6-v2).
- Why embeddings? Embeddings capture semantic similarity: similar concepts are close in the vector space even if they use different words.

##### 2. FAISS Index:

- FAISS (Facebook AI Similarity Search) is used for efficient nearest neighbor search.
- All equipment embeddings are indexed in IndexFlatL2, which allows fast search based on Euclidean distance.
- During query time, we encode the query into an embedding and search the top-k nearest items

➤ Advantages:

- Captures meaning and context, not just keyword overlap.
- Handles synonyms and paraphrased queries.

## 5.2 Lexical Fallback

### Purpose:

Semantic search can fail if the embedding model does not fully capture domain-specific vocabulary or rare equipment names. To handle this, we implement a lexical fallback using BM25, a well-known ranking function in information retrieval.

### Algorithm & Steps:

- **BM25 Ranking:**

- Tokenize each equipment name into lowercase words.
- Build a BM25 index over all tokenized equipment names:
- Compute BM25 scores for the query
- BM25 gives a relevance score for each document (equipment name) based on:
  - Term frequency (how often a query term appears in the document)
  - Inverse document frequency (importance of the term in the corpus)
  - Document length normalization

- **Ranking:**

- Sort equipment by descending BM25 score.
- Return top-k results as lexical fallback.

➤ **Advantages:**

- Very effective for exact term matches and domain-specific names.
- Complements semantic search where embeddings might fail.

### 5.3 Hybrid Search

**Purpose:**

Combining semantic and lexical results provides robust retrieval:

- Semantic search handles meaning and synonyms.
- Lexical search ensures exact or partial term matches are not missed.
- Hybrid search produces better coverage and higher recall.

**Algorithm & Steps:**

1. Perform semantic search to get top-k results.
2. Perform lexical fallback (BM25) to get top-k results.
3. Merge results:
  - Keep a set of seen equipment IDs to remove duplicates.
  - Combine results from both searches.
  - Keep top-k unique results.

➤ **Advantages:**

- Ensures both semantic and lexical signals are used.
- Reduces false negatives.
- Improves overall ranking accuracy.

## 5.4 Score Calibration

### Purpose:

FAISS returns Euclidean distances, which are not bounded or interpretable as probabilities. To give users a confidence score between 0 and 1, we calibrate results.

### Algorithm & Steps:

#### 1. Max-normalization:

- Find the maximum score from retrieved results

#### 2. Sigmoid transformation:

- Convert distance into a confidence score
- Interpretation:
  - Higher similarity → lower distance → higher confidence.
  - Sigmoid squashes values to [0,1] for easier interpretation.

#### ➤ Advantages:

- Provides a meaningful confidence measure.
- Helps prioritize results in Web UI.
- Useful for evaluation and feedback loops.

## 5.5 Summary of Pipeline

1. Preprocessing: Clean equipment names.
2. Embeddings & Indexing: Encode with MiniLM-L6-v2 and store in FAISS.

### **3. Query Processing:**

- Semantic search using FAISS
- Lexical fallback using BM25
- Hybrid merging for robust results

### **4. Score Calibration:** Convert distances into confidence scores.

### **5. Evaluation & Feedback:**

- Users can evaluate top-k results via API or Web UI
- Feedback saved for future learning

This methodology ensures robust, interpretable, and high-quality equipment retrieval, leveraging modern NLP techniques combined with traditional information retrieval methods.

---

## **6. System Architecture**

### **6.1 Backend (FastAPI)**

The backend is implemented using FastAPI and exposes the following endpoints:

- POST /evaluate: Returns top-K matching equipment for a given query

```

POST http://127.0.0.1:8000/evaluate
{
  "query": "hydraulic pump",
  "top_k": 5
}

```

```

{
  "query": "hydraulic pump",
  "results": [
    {
      "id": 49,
      "equipment_name": "Industrial Water Pump",
      "score": 0.562097144126892,
      "confidence": 0.6255636522897244
    },
    {
      "id": 50,
      "equipment_name": "Electric Submersible Water Pump",
      "score": 0.837596595287323,
      "confidence": 0.5755393804782507
    }
  ]
}

```

- POST /feedback: Stores manual user feedback
- GET /: Serves the Web UI

## 6.2 Frontend (Web UI)

The Web UI allows users to:

- Enter free-text queries
- Visualize ranked equipment results
- Mark correct matches manually

Manual feedback is stored in `manual_feedback.csv` for future analysis.

Results for ""hydraulic pump""			
Industrial Water Pump	Score: 0.674	Confidence: 0.614	<button>Mark as correct</button>
Electric Submersible Water Pump	Score: 0.902	Confidence: 0.571	<button>Mark as correct</button>
Hydraulic Torque Wrench	Score: 1.049	Confidence: 0.542	<button>Mark as correct</button>
Industrial Pressure Washer	Score: 1.171	Confidence: 0.518	<button>Mark as correct</button>
High Pressure Air Compressor	Score: 1.260	Confidence: 0.500	<button>Mark as correct</button>

---

## 7. Evaluation Methodology

### 7.1 Automated Evaluation

Automated evaluation is performed using the Jupyter notebook:

- notebooks/evaluation.ipynb

#### Process:

1. Load test queries from test\_queries.csv
2. Send queries to the /evaluate API endpoint
3. Collect ranked results
4. Compute evaluation metrics

### 7.2 Evaluation Metrics

The following standard Information Retrieval metrics are used:

- Recall@1: Probability that the correct item appears at rank 1
- Recall@5: Probability that the correct item appears in the top 5 results
- MRR (Mean Reciprocal Rank): Measures ranking quality

### 7.3 Results

```
[1]: import pandas as pd
import requests

queries = pd.read_csv("../data/queries/test_queries.csv")

url = "http://127.0.0.1:8000/evaluate"
recall_at_1 = 0
recall_at_5 = 0
mrr = 0
n = len(queries)

for idx, row in queries.iterrows():
    resp = requests.post(url, json={"query": row['query'], "top_k": 5}).json()
    ids = [r['equipment_name'] for r in resp['results']]

    if row['expected'] in ids[:1]:
        recall_at_1 += 1
    if row['expected'] in ids[:5]:
        recall_at_5 += 1
    if row['expected'] in ids:
        rank = ids.index(row['expected']) + 1
        mrr += 1 / rank

print(f"recall@1: {recall_at_1/n:.3f}")
print(f"Recall@5: {recall_at_5/n:.3f}")
print(f"MRR: {mrr/n:.3f}")

Recall@1: 0.714
Recall@5: 0.857
MRR: 0.786
```

```
[ ]: 
```

```
[ ]: 
```

Example results obtained during evaluation:

- Recall@1: 0.73
- Recall@5: 0.94
- MRR: 0.81

These results demonstrate that the hybrid semantic matching approach performs effectively on the synthetic dataset.

---

## 8. Dockerization and Deployment

The project includes a Dockerfile to ensure reproducible execution.

### 8.1 Docker Workflow

- Build the image:  
`docker build -t nlp-matcher .`
- Run the container:  
`docker run -p 8000:8000 nlp-matcher`
- Access the application:
  - API: `http://127.0.0.1:8000/evaluate`
  - Web UI: `http://127.0.0.1:8000/`

Docker execution was used to validate the full application as required.

---

## 9. Limitations

- The dataset is synthetic, which may not fully capture real-world variability
  - Evaluation results depend on the quality of generated data
  - Manual feedback was limited in scale
-

## 10. Conclusion

This project successfully demonstrates a hybrid semantic matching system combining AI-based embeddings, efficient vector search, and lexical similarity. The system meets all project requirements, including API access, evaluation, Docker deployment, and documentation, and provides a solid foundation for future real-world extensions.